

Software Interrupts

Irvine Edition IV : Section 13.1.4

Software Interrupts

- **Definition :** A **software interrupt** is a special **call** to an procedure previously defined as part of the Operating System
 - Alternate Terminology : TRAP, System Call
 - Implemented using a hardware mechanism:
interrupt service routine (ISR).
- **Examples :**
 - DOS Functions
 - Print a string message Functions
 - Exit
 - Character Input functions
 - Printer Output functions
 - BIOS Interrupts
 - Video Display
 - Disk I/O functions
 - Keyboard
 - Printers

Application Subroutines versus SW Interrupts

- Subroutines are part of an **application** program.
 - Developed during program development : assemble, link and load.
 - During assembly, the *assembler* determines the target's address offset
 - During linking, the application program is *linked* with other software that has also been assembled.
 - During loading, the segment values for the application program are initialized (May be different each time)
 - Whenever a change is made any file in the application program and/or a software library, the entire application must be re-built :
 - It must be assembled, linked and loaded again. Why ?

Application Subroutines versus SW Interrupts

- Traps **are like** subroutines
 - Transfer of control to an encapsulated activity terminated by a return to the invocation point.
- Traps **are different than** subroutines
 - They are not part of the application program
 - They are not part of the program development process
 - The OS is not a software library.
 - Your program's OBJ file are not linked with OS OBJ files.
- Key Difference : The O/S is a permanent resident of memory whereas the application is temporary resident

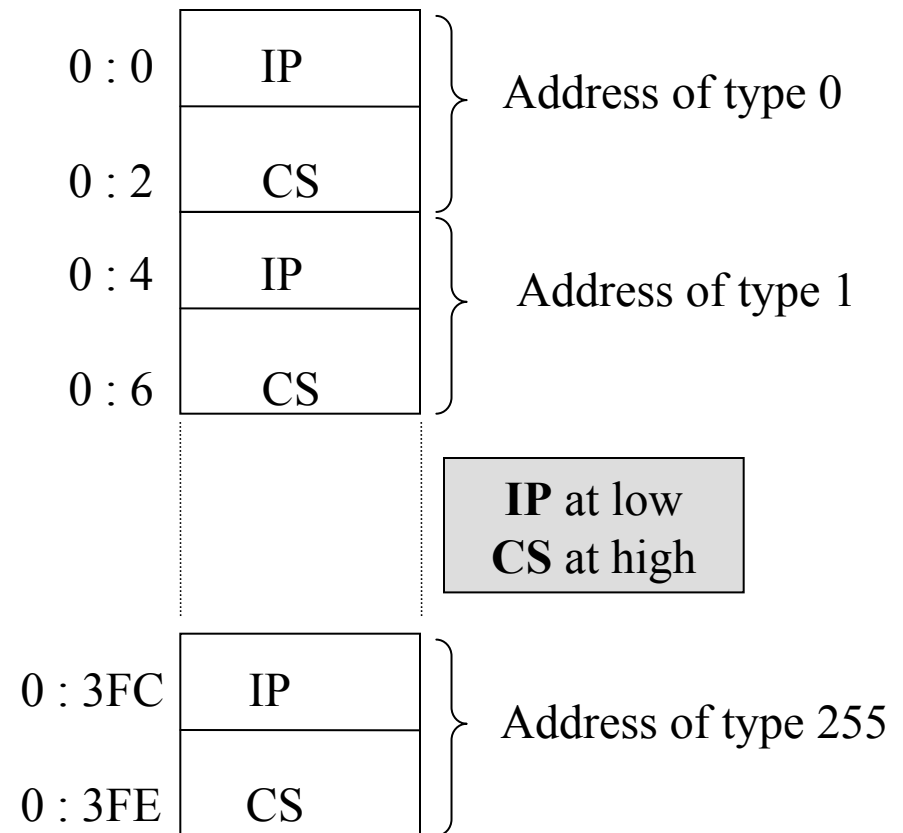
An application program is NOT LINKED with the OS

How does your program “call” the OS Procedure ?

- **First-Step Solution** : Locate each OS Procedure at “reserved” global locations
 - Locations are published as part of the OS (Contract/Policy)
 - Analogy : Direct memory addressing mode
 - What if the OS procedure is revised ?
- **Second-Step Solution** : Locate an array of addresses at a “reserved” global location containing the addresses of the OS procedures.
 - Analogy : Indirect Memory Addressing.
 - O/S puts pointers to O/S procedures in variables when O/S is loaded
 - Application programs are developed under the assumption that these global variables exist
 - What if the OS procedure is revised ?
- **Question** : Do we already have an array of addresses at a reserved location ?

8086 Vector Table

- An array of 256 entries located at the reserved memory location 0:0
 - Each entry is an address of an **interrupt service routine (ISR)**.
 - The address is a FAR Pointer (CS:IP pair) (32-bits = 4bytes)
 - The array occupies addresses from 0:0 to 0:3FF (256*4)
- Each entry is identified by unique "**interrupt-type**" (number), ranging from 0 to 255
 - If interrupt-type = i then the offset to relevant entry in the vector table = $0000H + 4 * i$



Software Interrupt (or TRAP) Instruction : INT

- Software interrupts are based on the stored pointer approach
 - Because new pointers can be installed, offer “dynamic subroutine invocation”

• Subroutines are identified by their label

• Subroutines are invoked with CALL instruction

• SW Interrupts are identified by their interrupt-type

• SW Interrupts are invoked with INT instruction

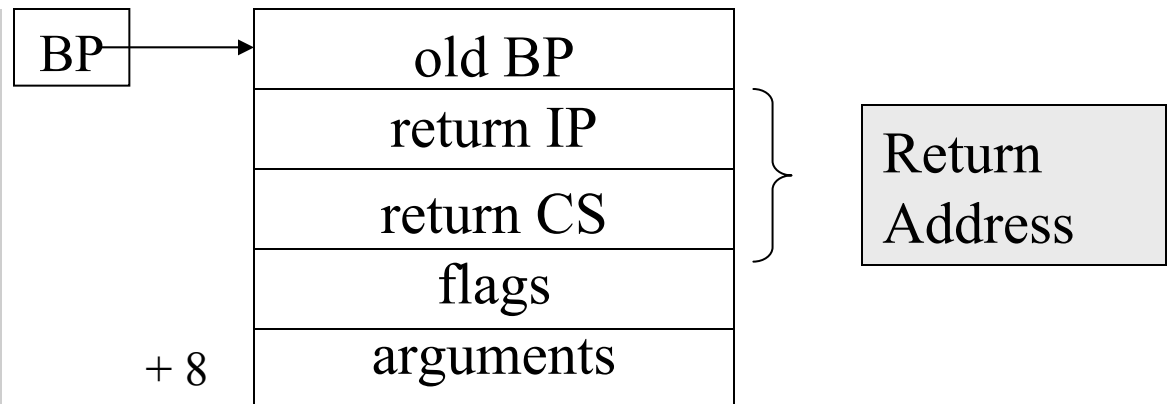
- 8086's INT Instruction
 - It is based on the globally assumed vector table at 0:0
 - Syntax : INT i where i = index into vector table (0..255)
- Example : INT 5
 - Invokes ISR whose FAR address is stored in 6th element of vector table at address = 0:14h (5*4=20)

Software versus Hardware Interrupts

- Software interrupts : invoked by execution of INT instruction
- Hardware interrupts : invoked by event on external I/O device
- **But both share the same execution semantics**
 - Push FLAGS register
 - Clear IF and TF bit in flags
 - Push CS and IP
 - Fetch new CS:IP from $0 : n*4 + 2$ and $0:n*4$
- Both thus share the same stack frame

Can we pass parameters to an ISR for a HW interrupt ? For a SW interrupt ?

- Is it different than passing parameters to subroutines ?



Subroutines versus Software Interrupts

- Both Subroutines and software interrupts transfer control to an encapsulated activity that is terminated by a return to the invocation point.

; subroutine initialization

none

...

; call set up

push arguments

CALL subr

ADD SP, 2*numArgs

...

subr:

standard entry code

access param's [BP + 4⁺]

standard exit code

RET

; ISR initialization

install address at 0:4*n

...

; call set up

push arguments

INT **n**

ADD SP, 2*numArgs

...

subr:

standard entry code

access param's [BP + **8**⁺]

standard exit code

IRET

n=0..255

We've already been using DOS Functions

What can you say about parameter passing to DOS Functions ?

```
.data
```

```
message db "Hello, world!", 0dh, 0ah, '$'
```

```
.code
```

```
; Print a string
```

```
    MOV     AH, 9
```

```
    MOV     DX, OFFSET message
```

```
    INT     21h
```

```
; Exit to DOS
```

```
    MOV     AX, 4C00h
```

```
; AH = 4Ch
```

```
; AL= 0 (exit status)
```

```
    INT     21h
```

INT 21h is the DOS Function Call

How does this one ISR handle all the services?

- The type of service is passed in as a parameter
- But parameters are passed by register, not the stack!!

AH=1 Keyboard input

AH=2 Character output

AH=5 Printer output

AH = 9 String output

AH = 4Ch Terminate program

- What about ISR return-value?
 - Depends on the service code parameter!

BIOS Function Calls

Other examples : BIOS Interrupts

- INT 10H for video display functions
- INT 13h for disk I/O functions
- INT 16h for keyboard functions
- INT 17h for printers functions

Intel 8086 Vector Table

- The 256 vector table has been mapped out – See Appendix C
- A *brief* look:

Interrupt Types

0 ... 1F

20h

21-24h

33h

60-6Bh

80-F0h

F1-FFh

Descriptions

Reserved by Intel

Includes : BIOS 10-16h

Terminate a COM program

DOS Functions

Mouse Functions

Available for Applications

Reserved

Available for Applications

The 5 Dedicated Interrupt (0..4)

- **Interrupt 0** (divide error)
 - Invoked **by the CPU** after a DIV or IDIV if the calculated quotient is larger than the destination
 - How big is the quotient if an attempt is made to divide by 0?
- **Interrupt 1** (single step)
 - Used by debuggers to support single stepping
 - If TF flag set, **CPU invokes** this ISR after executing most instructions
 - TF is cleared as part of the INT execution (after the flags are pushed)
 - Why is TF cleared ?
 - When ISR starts executing, processor is no longer in single-step mode
 - It avoids an infinite loop!

The 5 Dedicated Interrupt (0..4)

- **Interrupt 2** (non-maskable interrupt)
 - **Hardware** interrupt which cannot be disabled. More later....
- **Interrupt 3** (breakpoint interrupt)
 - A special version of the INT instruction that is encoded in one byte: CEH
 - Used to provide breakpoint capabilities for debuggers
- **Interrupt 4** (overflow interrupt) **INTO**
 - If OF set when INTO instruction is executed, the **CPU invokes** this ISR
 - Used in numeric libraries to trap overflow errors
- Higher processors (80186, 80286, etc.) have additional dedicated interrupts
 - IBM / Microsoft decided to use interrupts reserved by Intel for their own purposes. It caused problems when the AT was released